

AgentMPI: A Message Passing Interface for Multi-Agent Systems

AgentMPI Project

ABSTRACT

Multi-agent LLM systems are in the position parallel computing occupied around 1991: every group writes its own coordination layer, none of them compose, and there is no portable interface one writes a *harness* against. We argue that the interface parallel computing converged on—MPI [34, 35]—transplants, and we specify, implement and evaluate AGENTMPI: a protocol for writing multi-agent systems, not a multi-agent system.

The transplant is not mechanical. Five asymmetries between an MPI process and an LLM executor invert several of MPI’s answers. The scarce resource is the receiver’s *context window* rather than bandwidth, so MPI’s eager limit becomes a token-denominated flow-control problem and MPI’s “safe program” discipline becomes a mechanically checkable property we call *context safety*. The reduction operator costs seconds rather than nanoseconds, so algorithm selection must minimise operator applications on the critical path rather than message volume: at $p = 128$ the schedule MPI prefers for short messages performs 896 operator applications where reduce-then-broadcast performs 127. And because a shared control plane replaces a point-to-point network, MPI’s barrier hierarchy collapses entirely.

We also report a limitation of agent-evaluated reductions that we found rather than predicted, and a mechanism for it. A canonical reduction tree makes an agent reduction *reproducible*; it does not make it *consistent*, because two branches can meet the same conflict and resolve it oppositely with no node in a position to notice. We formalise this as the *locality of merge*, and introduce *conflict lifting*: enlarging the operator’s codomain with a conflict set whose join is a semilattice, so the conflicts reaching the root are identical for every tree shape and the root arbitrates each exactly once.

AGENTMPI ships as a specification, a reference runtime organised around a six-operation abstract device interface with 3 independent transports, and a conformance suite of 279 protocol assertions and 117 device assertions run unchanged against every transport. We evaluate with real agent populations: paired eight-rank translation runs, a hundred-rank run under executor oversubscription, and deterministic microbenchmarks that fit the cost model. Two results are negative and reported as such.

KEYWORDS

multi-agent systems, message passing, MPI, collective communication, fault tolerance, LLM agents

1 INTRODUCTION

A multi-agent LLM system today is written the way a parallel program was written in 1991. The author picks a coordination pattern—a group chat, a supervisor tree, a directed graph of handoffs—and implements it from scratch against whatever agent SDK is at hand. The result does not compose with anyone else’s, it cannot be moved to a different agent host without rewriting, and its failure modes are rediscovered independently by every group that builds one.

Parallel computing left that state by agreeing on an interface [34, 44]. The MPI Forum did not standardise a runtime, a language, or a programming model; it standardised the *calls a program makes* to communicate, and left the process manager, the algorithms and the transport to implementations. That decision is why a program written in 1996 still runs, and why a dozen vendors shipped MPI-1 within two years of the document.

This paper asks whether the same move is available now, and answers that it mostly is. AGENTMPI is a protocol you write a harness *with*, in the sense that MPI is a library you write a parallel program *with*. It does not decide how many agents to run, what they should do, how to prompt them, or how to recover; it provides the mechanisms with which those decisions are expressed. It is deliberately semantics-thin—a message is an opaque body plus an envelope—which is a considered rejection of the KQML/FIPA-ACL lineage whose mentalistic semantics proved unverifiable, in favour of MPI’s stance: standardise the mechanism, leave meaning to the application.

The transplant is not mechanical. Five properties of an LLM executor differ from an MPI process (§3), and each forces a deviation. Two of MPI’s algorithm-selection rules invert. One of MPI’s cost terms, the reduction operator, goes from free to dominant. One of MPI’s non-goals, fault tolerance, becomes the common case. And one resource with no MPI analogue—the receiver’s context window—turns algorithm selection from an optimisation into a feasibility test.

Contributions.

- (1) A specification (§4) that maps MPI’s chapters onto agent coordination and states, for each, whether the mechanism transfers, transfers with a changed unit, or does not transfer. It is written in the standard’s own idiom: normative text with explicit rationale.
- (2) *The locality of merge* (§5), a limitation of agent-evaluated reductions that we observed in a real run and could not previously name, together with two mechanisms—conflict lifting and invariant verification—and a shape-independence property for the first.
- (3) A cost model and re-derived selection rules (§6) in which the operator term dominates, with the resulting inversions stated and measured.
- (4) A reference implementation organised around a narrow waist (§7): six device operations, 3 independent transports, and one conformance suite run against all of them. The suite is the artifact that makes this a specification rather than a library, and it found defects that only exist on one transport.
- (5) An evaluation with real agent populations (§8), including two negative results we report because they are the informative ones.

Table 1: The five asymmetries that drive every deviation.

	MPI process	AGENTMPI executor
A1 Determinism	deterministic	may produce a different, plausible, wrong result
A2 Scarce resource	bandwidth	context window
A3 Operator cost	nanoseconds; free	seconds to minutes; dominant
A4 Latency	tight	heavy-tailed; max of p samples
A5 Failure	fail-stop, rare, fatal	frequent, partial, sometimes silent

2 BACKGROUND: WHAT MPI STANDARDISED

MPI was assembled in 1992–1994 [14, 34, 44] from a field of incompatible message-passing libraries—p4, PARMACS, Chameleon [24], Zipcode, PVM, Express, NX—each of which solved part of the problem and none of which was portable. The Forum’s rules were unusually restrictive and unusually effective: standardise existing practice rather than invent; decline everything not yet understood.

MPI-1’s *non-goals* were as consequential as its goals. It is not a language. It does not provide shared memory, debugging, I/O, threads, task management, or fault tolerance. Declining those kept it small enough to implement.

Three of its decisions matter here.

Communicators. A communicator pairs a group with a private *context* that partitions the message namespace. Pre-MPI libraries addressed messages by destination and tag, so a library’s messages could be intercepted by its caller whenever they collided on an integer. MPI’s designers retrospectively identify the context as the reason MPI could support libraries at all [23, 27, 41]. The agent-world version of that collision is not hypothetical: a reviewer’s critique intended for the author of module A being consumed by the author of module B is the same bug, and the “one shared group chat” architecture makes it certain rather than merely possible.

Safe programs. MPI deliberately refuses to guarantee buffering. A program that completes only because an implementation happened to buffer is called *unsafe*, and the standard’s advice is to test a program by replacing every standard-mode send with a synchronous one. The test is mechanical and almost nobody runs it, because the penalty is a deadlock on someone else’s machine.

Performance portability. The standard specifies semantics; implementations choose algorithms. The Hockney model $T = \alpha + \beta n$ and its successors [3, 13] let an implementation pick a collective algorithm by message size and process count [38, 42]. Evidence on which parts of MPI applications actually use [32] is what motivates our conformance levels.

3 WHY AN AGENT IS NOT A PROCESS

Table 1 lists them. Their consequences run through the rest of the paper: A2 forces the eager threshold to be denominated in tokens (§4.2); A3 inverts algorithm selection (§6); A4 motivates quorum collectives; A5 requires the whole fault-tolerance chapter;

A1 requires reduction *reproducibility* and, separately, *consistency* to be declared properties (§5).

One asymmetry runs the other way, and is easy to miss. In MPI the control plane is a point-to-point network. In AGENTMPI it can be, and in our implementation is, a *shared medium* every rank can read. That single structural difference collapses MPI’s barrier hierarchy: a counting barrier costs $2(p-1)$ operations against dissemination’s $p[\log_2 p] - 510$ against 2,048 at $p = 256$ —so it is cheaper at every size, and it is also the only algorithm that can name the ranks that did not arrive. MPI avoids counting barriers at scale for a reason that does not apply here.

4 DESIGN

We summarise the specification; the normative document accompanies the artifact.

4.1 Ranks, epochs, and identity

A rank is a *durable role*, not a process: a number, a mailbox, a context budget and an accumulated ledger. Its physical embodiment is an *executor* bound to it for a bounded interval, identified by an *epoch*. In MPI a rank *is* a process and the identification is harmless because MPI processes live for the whole job; an agent session does not, so tying rank identity to session identity would make every timeout a communicator rebuild.

The epoch is a fencing token. Leases alone are insufficient: an executor cannot know its lease expired mid-step, so between expiry and its next call two executors believe they own the rank. A monotone token checked on every operation closes that window and makes a zombie harmless rather than corrupting.

Identity is *ambient* but must be *assertable*. This is a correction we were forced into. Ambient identity eliminated the error we designed it to eliminate—agents passing the wrong rank—and replaced it with a worse one: agents silently *being* the wrong rank, because on one host shell sessions were shared between concurrent agents and the rank variable was overwritten between calls. The fix is not to abandon ambient identity but to make the assertion cheap.

4.2 Flow control in tokens

Every rank has a context budget; delivering a body charges it. Below an *eager threshold* (700 tokens by default) a body is pushed into the receiver; above it, the receiver gets an envelope and a content-addressed handle and materialises explicitly. This is MPI’s eager limit with the unit changed from bytes to tokens, and it exists for the same reason: below the limit, pushing unsolicited is cheaper than a handshake; above it, the receiver’s buffer—here, its attention—is too precious to fill without permission.

Long contexts are not free even when they fit: model accuracy degrades with position and length [33], and serving systems treat the window as a managed resource [29]. Both are reasons to account for it in the protocol rather than in the prompt.

Two things do not transfer. First, in MPI eager and rendezvous are invisible because both deliver the same bytes; here they differ in *what arrives*, so they must be visible. Second, exceeding the buffer does not deadlock, it silently degrades the receiver’s output.

That second point makes MPI’s safe-program discipline more important here, not less, so we mechanise it. Define a program

context-safe if it completes for every assignment of unexpected-message budgets, however small. The reference implementation runs a harness’s communication skeleton under zero-buffer semantics, reports the ranks that wedge and any cycle among them, and re-runs with every send declared rendezvous so it can say whether the harness is repairable by transport choice alone. A ring in which every rank sends before it receives—the shape agent harnesses write by default—is diagnosed with its cycle named and the repair stated.

4.3 Views and contracts

An MPI derived datatype is a declarative description of how to access a buffer, and its value is that the access pattern is data the library can optimise rather than control flow it cannot see. The scarce resource here is context, so the analogous question is not “which bytes” but “which tokens, and how few”. A *view* is a deterministic, model-free projection: `head:n, grep:pat, keys:a,b, outline, stat`. Semantic compression is deliberately *not* a view, because it costs an operator application the harness must be able to see.

Contracts transplant the second job of MPI datatypes: a type signature checked at both ends, turning silent corruption into a loud error. A contract may also require a payload to identify itself, which turns a misrouted scatter slice into an error at the receiver rather than a plausible wrong answer three phases later.

4.4 Collectives

Labelled, not positional. MPI identifies the *k*th collective by program order, which is safe when the caller is a compiled program. An executor’s program order is not: it retries, skips, and reorders. Positional identification would pair one rank’s second call with everyone else’s first and return a confidently wrong result. Labels cannot make a reordered schedule complete—nothing can—but they make the mismatch *named*.

Gather returns a *manifest*, not a concatenation. At $p = 128$ with 4,000-token contributions, an inlining allgather charges one rank 508,000 tokens and moves 65,024,000 in total; a handle-based one charges 5,080 and moves 650,240, a factor of 100 in peak residency. Concatenation is not a reasonable default at any interesting p .

Barriers carry a *policy*—wait, raise, proceed, shrink, revoke—because only the harness author knows whether a missing participant degrades the result or kills it. Collectives may carry a *quorum*, which *releases* without *closing*: a straggler still passes through, because a quorum that closed would guarantee that precisely the slowest ranks fail.

4.5 Windows

Every multi-agent system grows a shared scratchpad, and every one gets the concurrency wrong the same way: two agents read it, both edit a private copy, and the second write discards the first’s work. MPI-3’s one-sided chapter supplies both the vocabulary and the tools. `put` is the blind overwrite that loses work; `accumulate` combines atomically, so “union this finding into the shared findings” needs no lock; `compare_and_swap` is how work is claimed, and unlike a lock it cannot be held by a dead executor.

Two departures. Locks are *leased* and carry a fencing token, because an executor can wander off where an MPI process cannot.

And enumeration is free while reading is not: listing a window’s keys with their sizes and writers costs a few tokens per key, which is what makes a blackboard usable by an executor with a bounded context.

4.6 Fault tolerance

ULFM’s stance [5, 6] adopted wholesale: expose failure, provide mechanisms not policy. `revoke` is the non-obvious primitive—when a rank fails the *survivors* are the problem, blocked inside collectives that can never complete—and revocation makes them all fail fast so they reach the recovery path together. `shrink` derives a communicator over an *agreed* survivor set, and also offers FT-MPI’s [16] in-place mode, because renumbering invalidates a rank-to-work mapping that for agents is baked into prompts.

Recovery is by *briefing*, not checkpoint. There is no memory image to restore, so process checkpointing has no analogue; what does is durable-execution replay. The dominant failure here is a confident wrong answer, which is the exact analogue of silent data corruption, and HPC’s answer to that is not checkpoint/restart—which faithfully preserves the corruption—but algorithm-based fault tolerance [8, 28]. A replacement is given the answers to five questions it cannot guess: what was I assigned, what did I publish, what did I receive, what did I promise that is outstanding, and what did I record for myself.

Detection is lazy, local, and two-phase; it is an eventually perfect detector in Chandra and Toueg’s classification [12], and the leases carry fencing tokens for the reason Gray and Cheriton give [20]. A single-phase lease detector convicted 1091 times in twenty minutes on a real host, because a thinking executor and a dead one look identical. A blocked rank must renew its own lease while waiting; omitting that once made a barrier convict everyone who arrived early.

5 THE LOCALITY OF MERGE

MPI requires a user reduction operator to be associative and then reorders freely. The guarantee is deliberately weak—floating-point `MPI_SUM` is not bitwise reproducible across process counts [2]—and practitioners accept it because the discrepancy is bounded by rounding error.

An operator implemented by a language model is not associative, is lossy, is expensive, and is non-deterministic, and the discrepancy is bounded by nothing. AGENTMPI therefore makes the algebra *declared* rather than assumed and refuses schedules the declaration does not license. That much is straightforward.

What is not straightforward is a failure we found rather than predicted.

In an eight-rank reduction over interface proposals, two branches of the tree independently encountered the *same* conflict—which module defines the shared exception type—and resolved it in *opposite* directions. Both resolutions were recorded, so the merged result contained two contradictory rulings under one identifier, and the tree had no way to notice: each merge saw a locally consistent pair. Two ranks then implemented different halves of the result. Separately, the same reduction dropped

four of eight modules from one section of its output despite an explicit instruction never to drop a module.

A canonical tree shape makes a reduction *reproducible*. It does not make it *consistent*, and conflating the two is the mistake.

The observation, stated generally. Let a reduction be a fold of a binary operator $\oplus$ over a tree T whose leaves are the ranks’ contributions, and call a predicate φ a *global invariant* if its truth depends on the whole leaf multiset. An internal node of T sees exactly two operands. If two nodes u and v , neither an ancestor of the other, both encounter evidence bearing on φ , then no node of T is in a position to reconcile them: their lowest common ancestor receives u ’s and v ’s *outputs*, and the information that would reveal the inconsistency—what each decided, and why—is not in $\oplus$ ’s codomain. Enlarging the tree cannot help. Only enlarging the codomain can.

Conflict lifting. Let the operator’s codomain be $V \times C$, where C is a set of undecided conflicts, and require that an operator which cannot decide a contested item from its two operands alone *lift* it into C rather than decide it. Let C ’s join be set union.

PROPOSITION 5.1 (SHAPE INDEPENDENCE). *The conflict set arriving at the root of a fold over $\oplus$ is the same for every tree shape over the same leaf multiset, and no contested item is decided more than once.*

The argument is two lines. The conflict component’s join is associative, commutative and idempotent, so it is a semilattice fold, and a semilattice fold is independent of the order and shape of the fold. The value component never decides a contested item, so it cannot decide one twice. Arbitration at the root therefore decides each conflict exactly once, over the same set, whatever the schedule.

We verify this exhaustively: over every binary fold order at $p \leq 6$ (all Catalan($p - 1$) shapes) and over random orders to $p = 32$, the conflict set is invariant, while a control operator that decides locally is not.

Why it is cheap. Lifting buys the “one ruling per key” invariant for one extra operator application and $O(|C|)$ tokens. The alternative that also buys it—a serial chain, in which the accumulator carries every prior decision—costs $p - 1$ critical-path applications against a tree’s $\lceil \log_2 p \rceil$: at $p = 64$, 63 against 6.

Implementation warning. Represent the lifted state as a map from key to the *set* of distinct values seen, and derive the value/conflict presentation from it. The obvious implementation—lift on first disagreement and remove the key from the value—is *not* idempotent: a third contributor finds the key absent and reinstates it, so the conflict set depends on fold order after all. We shipped that version, and its own shape-independence test caught it.

The other half. “Never drop a module” is not a conflict any pair of operands exhibits—every local merge is individually faithful—so lifting is blind to it. It is a property of the leaf multiset and the result, so an operator may declare a *global invariant* that is checked once, after the reduction closes, and reported with the missing items named.

Table 2: Total operator applications, allreduce. Recursive doubling and reduce-then-broadcast put the same number on the critical path.

p	8	16	32	64	128
recursive doubling	24	64	160	384	896
reduce + broadcast	7	15	31	63	127
ratio	3.4	4.3	5.2	6.1	7.1

Neither mechanism makes an agent operator correct. They make two specific, observed, expensive failure modes into loud errors, which is the most a protocol can honestly claim.

6 COST MODEL AND ALGORITHM SELECTION

For an operation moving n tokens,

$$T = \alpha + \beta n + \gamma k + \lambda, \quad C = c_{\text{tok}} n + c_{\text{op}} K,$$

where γ is the cost of one operator application by an executor, k the number on the critical path, K the number in total, and λ the executor’s own heavy-tailed think time.

Measured on the reference implementation (§8.1), $\alpha = 0.730$ ms and $\beta = 0.480$ $\mu\text{s}/\text{token}$ on the SQLite transport, giving a half-bandwidth point of 1,521 tokens—above the size of a typical agent artifact, so latency-optimal algorithms win over a wider range of sizes than in MPI. Against them, γ for an LLM merge is seconds to minutes. At $p = 128$ the entire protocol cost of an allreduce is 15.09 ms, which is 0.215 % of the total against a one-second operator and 0.0072 % against a thirty-second one.

Two of MPI’s rules invert.

(a) *Runtime operators want no tree.* When the implementation can apply the operator—set union, JSON merge, max—a shared control plane folds every contribution in place: one round, zero messages. A tree adds rounds and buys nothing. This is the in-network-aggregation regime, and it is the common case for the operators harnesses *should* be using. Sweeping γ confirms it: at $\gamma = 0$ the winner is `flat`, and at $\gamma = 30$ s it is `reduce_bcast`, with the flat schedule $10 \times$ behind (1890.0 s against 180.0 s).

(b) *Agent operators want MPI’s tree, and reject MPI’s best tree.* A binomial tree puts $\lceil \log_2 p \rceil$ applications on the critical path against a chain’s $p - 1$; the catalogue also includes Bruck’s algorithm [9] and Rabenseifner’s [39], and the scan schedules follow Ladner and Fischer [31] and Bletloch [7]. But recursive-doubling allreduce—MPI’s standard choice for short messages [38, 42], precisely because redundant arithmetic is free—performs $p \lceil \log_2 p \rceil$ applications *in total* (Table 2) for the same critical path. The algorithm that wins for bytes loses badly for agents, and it loses on the *price* axis, which is why the model reports time and price separately.

Admissibility precedes optimisation. An algorithm whose peak per-rank residency exceeds a context budget is *infeasible*, not slow. At $p = 128$ with 4,000-token payloads, every inlining allgather in the catalogue is rejected with its peak residency named; the handle-based variants are admissible. MPI has no analogue: there, every algorithm runs.

7 IMPLEMENTATION

7.1 The narrow waist

MPICH’s portability comes from an abstract device interface separating portable MPI semantics from transport [10, 21, 22]: Open MPI makes the same move with MCA frameworks. AGENTMPI copies the structure, for a different reason: an interface with one implementation is a library, and what would make this a standard is that the semantics can be stated and *checked* independently of any transport.

Six operations sit at the waist: append (durable, totally ordered), match (atomic first-fit claim), scan, cas, lease, clock. Everything above—communicators, matching, collectives, windows, the ledger, revoke/shrink/agree—is written in terms of those six. Why not fewer: append and scan suffice for a single writer, but a receive must not be handed to two ranks, and that atomicity cannot be recovered above the device. Why not more: window history is a scan over versions, the failure detector is a scan over leases, and the deadlock detector is a scan over blocked receives.

Three transports ship, sharing no code below the interface: SQLite in WAL mode, a filesystem journal using one file per record with an advisory lock, and an in-process dictionary. 10,822 lines of implementation sit above the waist.

7.2 The conformance suite

117 device assertions and 279 protocol assertions, each naming the specification section it enforces, run unchanged against all 3 transports. The suite uses only public interfaces, so pointing it at a different implementation means changing one fixture.

It earns its keep. In its first pass it found six defects, of which two matter beyond this codebase. An administrative kill was retractable by its victim’s next heartbeat—which would have made every fault-injection measurement meaningless. And SQLite’s type affinity was returning integers as strings for four indexed fields, so a wildcard receive posted as `-1` read back as `"-1"` and silently stopped matching: a divergence that existed on *one* transport and would otherwise have surfaced as an inexplicable protocol bug months later. That is precisely what a suite parametrised over every transport is for.

7.3 The command binding

For an LLM executor the binding is a command-line tool, because that is the only interface an agent reliably has to a stateful library: it cannot hold a handle across turns, cannot link a shared object, and its “function calls” are invocations whose output lands in its context window. Eight properties are normative, each because its absence cost a run: assertable identity; an identity echo on every call; errors that carry a concrete next action and a retryable flag; internal retry; sizes in tokens; a free path to disk on every operation that hands back a payload; the manual as a command; and *print only commands that exist*. That last one has a test that walks every emitted command string through the real parser, because an early reduction directive printed a subcommand spelled with a space where the real one is hyphenated, and roughly ten agents reported it, several while peers were blocked behind them.

Table 3: E1, two paired arms at two sizes. Both metrics are null at both sizes, and the collective costs wall clock at both.

	$p = 8$		$p = 32$	
	gloss.	ctrl	gloss.	ctrl
Consistency	1.000	1.000	0.994	0.994
strict	0.9474	0.9474	0.9509	0.9632
Terms scored	10	10	15	15
Wall clock (s)	400	109	346	255
Result tokens	6,648	6,082	26,480	24,406

8 EVALUATION

8.1 Microbenchmarks

No model is involved, deliberately: the job is to measure the *protocol* so the agent experiments can attribute their time correctly.

Fitting $T = \alpha + \beta n$ by ping-pong regression over payloads from 1 to 12,800 tokens gives $\alpha = 0.730$ ms on SQLite, 3.947 ms on the filesystem journal, and 0.096 ms in memory. The half-bandwidth point is 1,521 tokens on SQLite and 8,104 on the journal.

One result was not anticipated: β agrees to within 1.5 % across all three transports. The per-token cost is serialisation and token counting *above* the waist, not transport below it. Only α is a property of the device—which is a useful thing for the cost model to know, and an argument that the waist is in the right place.

8.2 E1: translating a book

The task is native translation of a book, chosen because it is *almost* embarrassingly parallel. Each passage translates independently; what does not is the rendering of names and coined terms recurring across passages. Ignore the coupling and the Tin Woodman acquires four names; serialise to avoid it and you pay p executor latencies.

We split *The Wonderful Wizard of Oz* into 100 passages and extract 21 recurring terms by capitalisation frequency, filtered by a model-free rule (a candidate that also appears lower-cased anywhere in the corpus is dropped). The harness has six phases, each one collective: scatter the passages; an agent phase proposing renderings; an allreduce with union, which lifts disagreements rather than deciding them; arbitration at the root; a broadcast of the binding glossary; an agent phase translating; a gather.

The protocol is in the harness, not the prompt. An executor’s entire obligation is to read a prompt file, write a result file, and submit; the worker bootstrap prompt mentions no communicator, collective, barrier or glossary. That is the design claim stated as an artifact, and it is why the same experiment runs against a deterministic function or a hundred live agents without changing a line.

Scoring is mechanical. For each recurring term, we find which rendering each passage’s translation used and take the modal rendering’s share, weighted by how many passages contain the term. The candidate renderings searched for are pooled across *all* arms compared, which matters: our first scorer built the vocabulary from each run’s own term sheets, and since only the treatment arm produces term sheets, the control was scored against an almost-empty vocabulary and flattered by it.

The mechanism worked, and scaled as predicted. At $p = 8$ the reduction agreed 17 terms unanimously and lifted 3 as conflicts; at $p = 32$ it agreed 8 and lifted 12. Quadrupling the population quadrupled the disagreement it found, which is what one would expect and what the conflict machinery exists to surface. The root arbitrated each conflict once, at both sizes.

Every one of the 12 conflicts was the same kind: whether the Spanish definite article belongs in the rendering—*el Espantapájaros* against *Espantapájaros*, *la Ciudad Esmeralda* against *Ciudad Esmeralda*. Not one was a disagreement about the noun.

The result is null, at both sizes. Table 3. Weighted terminology consistency is 1.000 in all four arms; under the strict metric that counts the article as part of the decision, the control arm at $p = 32$ scores marginally *higher* than the treatment (0.9632 against 0.9509). The glossary cost $3.7 \times$ the control’s wall clock at $p = 8$ and $1.4 \times$ at $p = 32$.

We had predicted that the null at $p = 8$ was a scale effect and that a larger population would surface terms with weaker priors. The prediction was half right and instructive in the half that was wrong. The population’s “stated” conventions did diverge more at scale, exactly as predicted—four times as many lifted conflicts. But that divergence did not reach the artifact, because the thing the translators disagreed about was the definite article, and in running prose the article is determined by the grammar of the sentence rather than by a lexical choice. Each translator wrote *el Espantapájaros* or *Espantapájaros* as the sentence required, whichever they had declared.

So the collective correctly identified a real disagreement in what the population said it would do, and that disagreement turned out not to matter for what the population produced. This is the mirror image of the caution in §5: agreement is not correctness, and here disagreement was not incorrectness either. The operational lesson for a harness author is the one we did not expect to be reporting—before paying a collective to remove a coupling, measure whether the coupling is costing anything.

We note the threats. This is $n = 1$ per cell, one model family, one language pair, and one corpus whose recurring terms are famous enough that a strong model has seen them. A task whose terminology is genuinely novel—an internal codebase, a new API—is where we would expect the collective to pay, and we have not run it.

8.3 E2: fault tolerance

The question here is not how a model behaves under failure—E1 and the conformance suite cover that—but whether the recovery machinery does what §4 says when ranks die at the worst moments. That has a definite answer, and getting it means killing ranks at controlled points many times, which is exactly what one should not pay a language model to sit through. Scripted ranks, $p = 12$, five scenarios, each a specification claim stated as a measurement. All 5 of 5 hold.

The last row is the one worth reading carefully (Table 4). With one rank dying mid-claim, 0 tasks were done twice and 1 was left orphaned: claimed by a dead executor and never finished. That is exactly what the protocol promises and no more. Compare-and-swap guarantees a task is not done twice; finishing a dead claimant’s

Table 4: E2: each row is a claim from the specification, measured.

Claim	Measured
A rank killed before a collective must not hang it	3 killed, 9 contributors, omission recorded
Revocation frees a survivor blocked <i>inside</i> a collective	freed after 1.32 s with ERR_REVOKED
Concurrent shrinks converge	11 ranks shrinking, 1 communicator
A replacement can resume from the briefing alone	5/5 published cells, plus the open collective and the memo
Claimed work is never done twice	24 tasks, 23 done, 0 duplicated

task requires a supervisor, and reporting the orphan count is how a harness author learns how much they must still do themselves.

8.4 Oversubscription, and a launcher constraint

Our agent host caps concurrent sessions at ten. This is not an accident of our setup; every host we have used imposes some such cap, and it is exactly the condition that leaves a job with ranks nobody will ever start. The $p = 32$ runs above are served by 8 executors at oversubscription four.

The protocol’s answer is oversubscription, and it has a precise MPI analogue: `mpirun -np 100` on eight cores runs a hundred ranks on eight cores. It works here for the same reason it works there—a rank is a durable role whose state lives outside its executor—and it is only available because the protocol is in the harness. An executor serving ten ranks *agent-side* would have to be inside ten barriers at once and would deadlock immediately; here the harness thread blocks in the collective and the executor blocks on nothing.

The specification’s other answer is the *join deadline*: a rank’s lease starts when the rank is *requested*, not when its executor first calls in. Without it, a rank whose executor never starts has no lease to expire, so it is neither alive nor failed and every peer waits for it forever. With it, a launcher that can start six of twenty-two ranks produces sixteen detectable no-shows.

8.5 Two identities, and the one we thought was checked

Two identities travel with every task in these experiments. The *protocol identity* is the rank an operation acts as, which the binding lets a caller assert with `-expect-rank`. The *provenance label* is an environment variable we added so that a scale claim would have per-executor evidence; it is outside the protocol and nothing checks it.

On a host where shell sessions are shared between concurrently running agents, both drifted. Across 199 tasks and 19 executors the provenance label was wrong on 18 of them (9.0 %). In every run with one executor per rank, and in the $p = 32$ runs at oversubscription four, it drifted 0 times.

We first reported that the protocol identity held because it was asserted. That was wrong, and two executors told us so. One wrote: “the next call carried `-expect-rank 3` and still returned rank 2’s identity.” Another diagnosed it exactly: “the guard didn’t fire because the clobber replaced `-rank` and `-expect-rank together`.” Both

are right, for two independent reasons, and the reasons are more useful than the claim we had made.

The implementation defect. The assertion was wired into the library API and into the ordinary CLI operations, and *not* into the worker subcommand — the only surface any executor in these experiments ever touched. The conformance suite covered the paths we had thought about. What actually caught the one misroute that was caught was a different and weaker mechanism: the broker’s check that a submitted task belongs to a rank the caller serves. That check is a serve-set membership test, not an identity assertion, and it fires only at submit, after the work is done.

The design defect, which is the interesting one. An assertion compares two values the caller supplies. A corruption that rewrites *both* the acting rank and the assertion leaves them consistent, so the check passes and the operation acts as somebody else. No amount of asserting fixes this, because the assertion is made of the same material as the thing it asserts about.

The mechanism that does fix it was already in the specification and already implemented: a per-rank *launch token*, issued by the launcher and recorded in the job, which cannot be made self-consistent by an accident of the environment. We had the mechanism and not the property, because the path that needed it did not call it.

What the executors did instead. Both misrouted claims were handled correctly, by agent judgement rather than by the protocol. One executor abandoned its task with an accurate reason. The other declined even to abandon it, reasoning that give-up would also have acted as the wrong rank and would have discarded another executor’s work, and let the claim lease-expire instead — which the broker requeued, and its rightful server completed. That is a better decision than the protocol was in a position to enforce, and we would rather not have to rely on it.

All four defects are fixed and regression-tested: identity is asserted on the worker path, the emitted submit and give-up commands carry `-expect-rank` rather than leaving it to the agent to remember, and the identity flags are accepted on either side of the subcommand. The last of those is its own lesson (§9).

9 WHAT ONLY REAL AGENTS EXPOSED

Several rules in the specification exist because a run failed. They are worth listing separately, because they are the part of this work that could not have been derived from MPI.

- (1) **Executors give up far earlier than instructed.** One told to retry up to twenty times gave up after two and stalled its reduction tree. A protocol that depends on an executor’s persistence is not a protocol; the binding retries internally and the default is “keep waiting”.
- (2) **Blocking is not evidence of death.** A rank waiting inside a barrier made no runtime calls, so the detector convicted it for waiting, and every rank that arrived early was declared failed. A blocked rank must renew its own lease.
- (3) **Program order is not a collective identifier.** Hence labels.
- (4) **Ambient identity without assertion is unsafe.** Shared shell sessions made one rank accumulate nine initialisations, one from nearly every other agent.
- (5) **Never truncate a reduction operand.** A result may be summarised because its consumer is a reader; an operand may not, because its consumer is a function. Agents handed JSON clipped mid-string invented recovery hacks, one prefix-matching the content-addressed store by hand.
- (6) **Print only commands that exist** — including the bootstrap prompt, which is also a command the binding prints. Ours placed the identity flags before the subcommand, where a reader expects a global option and where our parser rejected them, and roughly thirty executors each lost their first call to `invalid choice`. Our mechanical check walked the strings the runtime emits and not the prompt every executor reads first. And any guard left out of a printed command is a guard that is present only when the agent thinks of it: the emitted `submit` omitted `-expect-rank`, and exactly the executors who noticed had it.
- (7) **Give every payload a free path to disk.** Agents that lacked one reached into the object store directly rather than pay to see what was already a file.
- (8) **A rank that never starts must still be detectable.**
- (9) **A conditional send with an unconditional receive is a deadlock.** The protocol can only bound the damage; the advice is to send unconditionally, possibly empty. The static checker catches it.
- (10) **When a protocol omits a mechanism, its users build it worse.** Eight agents given no coordination phase rebuilt an allgather by hand within ninety seconds of one another, then built runtime interface discovery on top at roughly five times the detection code and twice the total lines for the same graded behaviour, with two of the eight introducing defects inside their own detection logic. AGENTMPI/1.0 therefore specifies interface declaration *and* verification, because the arms showed they are complementary: declaration alone reproduces the inconsistency of §5, and verification alone pushes the whole cost into every consumer.

10 DISCUSSION AND LIMITATIONS

What we have not shown. We have not shown that AGENTMPI produces better answers than an ad-hoc harness. Our one controlled quality comparison is null (§8), and we report it as such. What we have shown is narrower and, we think, still worth having: that MPI’s mechanisms have well-defined agent analogues, that several of MPI’s answers change in stated and measurable ways, that the resulting interface can be specified precisely enough to be conformance tested, and that real agent populations run against it.

We hot-patched a running job, twice. An executor in the scale run reported an `ImportError` from a module another session was in the middle of fixing. Appendix D of our own specification warns about exactly this, and the warning was inadequate: it said to pin the runtime *version*, and the version string had not changed, because what changed was the code. The runtime now records a content hash of its own source at job creation and logs a diagnostic when a later caller’s differs — advisory rather than fatal, because a developer

iterating should not be locked out and a two-hour run should not die over a reflowed comment, but recorded, because a run whose runtime changed halfway through should say so in its own journal.

Statistical power. Every agent experiment here is $n = 1$ per arm. Agent runs are expensive and not reproducible, which is also why tracing is unconditional and why we commit raw per-rank artifacts rather than aggregates.

Scale of the agent runs. Our largest *completed* agent run is $p = 32$ over 8 executors. We attempted $p = 100$ and it did not finish: the host’s session cap forced several waves of executors, and an operator error later destroyed the run’s journal. The pull queue handled the first problem well—a fresh wave could be added mid-run against exactly the ranks still queued, with no harness change and no restart, and 199 tasks were published across 19 executors—and the second was ours, not the protocol’s. We report the partial run’s identity measurement because it was committed before the journal was lost, and we do not report a hundred-rank result, because we do not have one. Microbenchmarks reach $p = 128$ with scripted ranks; the agent runs do not.

The strongest objection. MPI’s SPMD model presupposes a static rank set and interchangeable workers, and agent systems are neither. We think the objection is partly right: the fixed world size is a real constraint, and heterogeneous roles fit awkwardly into a model where rank r ’s program is the same as rank s ’s. Our defence is that a rank is a *role* rather than a worker, that communicators already express heterogeneity by partitioning, and that the constraint buys the thing the ad-hoc alternatives lack: a fixed membership against which a collective can be defined at all.

Agreement is not correctness. A protocol that makes agreement cheap makes it cheap to agree on something false. In our translation run a glossary bound a term to a rendering correct nearly everywhere, and a rank correctly obeying it produced a wrong sentence in the one passage with a different sense. A consistency metric scores a uniformly wrong decision as perfect.

11 RELATED WORK

Agent protocols. MCP [37] standardises client–server tool access, and A2A [1] and ACP standardise agent discovery and task hand-off. None provides a communicator, a collective, shared state with atomics, failure detection, or flow control; they are complementary rather than competing, and AGENTMPI could run over any of them as a transport.

Agent communication languages. KQML [17] and FIPA-ACL [18] define speech acts with mentalistic semantics [30]. The critique that those semantics are unverifiable from outside applies with more force to a language model, and we take MPI’s road instead.

Frameworks. AutoGen [45], LangGraph, CrewAI, MetaGPT and their relatives supply coordination *patterns*. Empirical work on how such systems fail [11] independently identifies several of the modes our failure model enumerates. They are the analogue of the pre-MPI libraries: each solves part of the problem, and none composes with another. The distinction we draw is that a framework makes the

coordination decisions and a protocol provides the mechanisms with which to make them.

Distributed computing. Win_fence is Valiant’s BSP superstep boundary [43]. Quorum collectives are stale-synchronous parallelism [26] expressed as a collective. Supervision with a bounded restart intensity is Erlang/OTP’s [4, 15]. Recovery by briefing is durable execution [36]. Windows are a blackboard in the sense of HEARSAY-II and Linda [19, 25], with MPI-3’s one-sided discipline imposed on it, and work claiming by compare-and-swap is the Contract Net’s descendant [40]. We claim the synthesis, not the parts.

12 CONCLUSION

MPI’s designers were solving a coordination problem under a particular cost model, and they were unusually explicit about which of their decisions followed from that model. That explicitness is what makes the transplant possible: one can take the mechanism, change the cost model, and see which answers change. Most of them survive. Two selection rules invert, one hierarchy collapses, one non-goal becomes the common case, and one resource with no analogue turns optimisation into feasibility.

The part that did not come from MPI is the part we would most like to be wrong about. A reduction evaluated by agents can be perfectly reproducible and still internally contradictory, because merges are local and invariants are not. We can name the problem, give it a mechanism whose correctness is a two-line semilattice argument, and check that mechanism exhaustively. We cannot make an agent operator correct, and we do not think a protocol should claim to.

REFERENCES

- [1] A2A Project. 2026. Agent2Agent (A2A) Protocol Specification, version 1.0.0. Linux Foundation / Agentic AI Foundation. <https://a2a-protocol.org/v1.0.0/specification/>
- [2] Peter Ahrens, James Demmel, and Hong Diep Nguyen. 2020. Algorithms for efficient reproducible floating point summation. *ACM Trans. Math. Software* 46, 3 (July 2020), 22:1–22:49. <https://doi.org/10.1145/3389360>
- [3] Albert Alexandrov, Mihai F. Ionescu, Klaus E. Schauer, and Chris Scheiman. 1995. LogGP: Incorporating long messages into the LogP model — one step closer towards a realistic model for parallel computation. In *Proceedings of the Seventh Annual ACM Symposium on Parallel Algorithms and Architectures (SPAA '95)*. ACM, 95–105. <https://doi.org/10.1145/215399.215427>
- [4] Joe Armstrong. 2003. *Making Reliable Distributed Systems in the Presence of Software Errors*. Ph. D. Dissertation. Royal Institute of Technology (KTH), Stockholm, Sweden.
- [5] Wesley Bland, Aurelien Bouteiller, Thomas Herault, George Bosilca, and Jack Dongarra. 2013. Post-Failure Recovery of MPI Communication Capability: Design and Rationale. *The International Journal of High Performance Computing Applications* 27, 3 (2013), 244–254. <https://doi.org/10.1177/1094342013488238>
- [6] Wesley Bland, Aurelien Bouteiller, Thomas Herault, Joshua Hursey, George Bosilca, and Jack J. Dongarra. 2012. An Evaluation of User-Level Failure Mitigation Support in MPI. In *Recent Advances in the Message Passing Interface (EuroMPI 2012) (Lecture Notes in Computer Science, Vol. 7490)*. Springer, 193–203. https://doi.org/10.1007/978-3-642-33518-1_24
- [7] Guy E. Blelloch. 1990. *Prefix sums and their applications*. Technical Report CMU-CS-90-190. School of Computer Science, Carnegie Mellon University.
- [8] George Bosilca, Rémi Delmas, Jack J. Dongarra, and Julien Langou. 2009. Algorithm-Based Fault Tolerance Applied to High Performance Computing. *J. Parallel and Distrib. Comput.* 69, 4 (2009), 410–416. <https://doi.org/10.1016/j.jpdc.2008.12.002>
- [9] Jehoshua Bruck, Ching-Tien Ho, Shlomo Kipnis, Eli Upfal, and Derrick Weathersby. 1997. Efficient algorithms for all-to-all communications in multiport message-passing systems. *IEEE Transactions on Parallel and Distributed Systems* 8, 11 (Nov. 1997), 1143–1156. <https://doi.org/10.1109/71.642949>
- [10] Darius Buntinas, Guillaume Mercier, and William Gropp. 2006. Design and Evaluation of Nemesis, a Scalable, Low-Latency, Message-Passing Communication Subsystem. In *Sixth IEEE International Symposium on Cluster Computing and the Grid*

- (CCGrid '06). IEEE, Singapore, 521–530. <https://doi.org/10.1109/CCGRID.2006.31>
- [11] Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. 2025. Why Do Multi-Agent LLM Systems Fail?. In *Advances in Neural Information Processing Systems 38 (NeurIPS 2025), Datasets and Benchmarks Track*. arXiv:2503.13657 <https://arxiv.org/abs/2503.13657>
 - [12] Tushar Deepak Chandra and Sam Toueg. 1996. Unreliable Failure Detectors for Reliable Distributed Systems. *J. ACM* 43, 2 (1996), 225–267. <https://doi.org/10.1145/226643.226647>
 - [13] David Culler, Richard Karp, David Patterson, Abhijit Sahay, Klaus Erik Schauer, Eunice Santos, Ramesh Subramonian, and Thorsten von Eicken. 1993. LogP: Towards a Realistic Model of Parallel Computation. In *Proceedings of the Fourth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP '93)*. ACM, 1–12. <https://doi.org/10.1145/155332.155333>
 - [14] Jack J. Dongarra, Rolf Hempel, Anthony J. G. Hey, and David W. Walker. 1993. *A Proposal for a User-Level, Message-Passing Interface in a Distributed Memory Environment*. Technical Report ORNL/TM-12231. Oak Ridge National Laboratory.
 - [15] Ericsson AB. 2025. Supervisor Behaviour — Erlang/OTP System Documentation. https://www.erlang.org/docs/29/system/sup_princ.html
 - [16] Graham E. Fagg and Jack J. Dongarra. 2000. FT-MPI: Fault Tolerant MPI, Supporting Dynamic Applications in a Dynamic World. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface (EuroPVM/MPI 2000) (Lecture Notes in Computer Science, Vol. 1908)*. Springer, 346–353. https://doi.org/10.1007/3-540-45255-9_47
 - [17] Tim Finin, Richard Fritzson, Don McKay, and Robin McEntire. 1994. KQML as an Agent Communication Language. In *Proceedings of the Third International Conference on Information and Knowledge Management (CIKM '94)*. ACM, Gaithersburg, Maryland, USA, 456–463. <https://doi.org/10.1145/191246.191322>
 - [18] Foundation for Intelligent Physical Agents. 2002. *FIPA ACL Message Structure Specification*. Technical Report SC00061G. Foundation for Intelligent Physical Agents. <http://www.fipa.org/specs/fipa00061/SC00061G.html>
 - [19] David Gelernter. 1985. Generative Communication in Linda. *ACM Transactions on Programming Languages and Systems* 7, 1 (Jan. 1985), 80–112. <https://doi.org/10.1145/2363.2433>
 - [20] Cary G. Gray and David R. Cheriton. 1989. Leases: An Efficient Fault-Tolerant Mechanism for Distributed File Cache Consistency. In *Proceedings of the Twelfth ACM Symposium on Operating Systems Principles (SOSP '89)*. ACM, 202–210. <https://doi.org/10.1145/74850.74870>
 - [21] William Gropp and Ewing Lusk. 1994. *An Abstract Device Definition to Support the Implementation of a High-Level Point-to-Point Message-Passing Interface*. Technical Report Preprint MCS-P342-1193. Mathematics and Computer Science Division, Argonne National Laboratory, Argonne, IL, USA.
 - [22] William Gropp, Ewing Lusk, Nathan Doss, and Anthony Skjellum. 1996. A High-Performance, Portable Implementation of the MPI Message Passing Interface Standard. *Parallel Comput.* 22, 6 (1996), 789–828. [https://doi.org/10.1016/0167-8191\(96\)00024-5](https://doi.org/10.1016/0167-8191(96)00024-5)
 - [23] William D. Gropp. 2001. Learning from the Success of MPI. In *High Performance Computing — HiPC 2001, 8th International Conference, Hyderabad, India, December 17–20, 2001 (Lecture Notes in Computer Science, Vol. 2228)*, Burkhard Monien, Viktor K. Prasanna, and Sriram Vajapeyam (Eds.). Springer, Berlin, Heidelberg, 81–92. https://doi.org/10.1007/3-540-45307-5_8
 - [24] William D. Gropp and Barry Smith. 1993. *Chameleon Parallel Programming Tools Users Manual*. Technical Report ANL-93/23. Argonne National Laboratory.
 - [25] Barbara Hayes-Roth. 1985. A Blackboard Architecture for Control. *Artificial Intelligence* 26, 3 (1985), 251–321. [https://doi.org/10.1016/0004-3702\(85\)90063-3](https://doi.org/10.1016/0004-3702(85)90063-3)
 - [26] Qirong Ho, James Cipar, Henggang Cui, Jin Kyu Kim, Seunghak Lee, Phillip B. Gibbons, Garth A. Gibson, Gregory R. Ganger, and Eric P. Xing. 2013. More Effective Distributed ML via a Stale Synchronous Parallel Parameter Server. In *Advances in Neural Information Processing Systems 26 (NIPS 2013)*. 1223–1231.
 - [27] Torsten Hoefer and Marc Snir. 2011. Writing Parallel Libraries with MPI — Common Practice, Issues, and Extensions. In *Recent Advances in the Message Passing Interface, 18th European MPI Users' Group Meeting, EuroMPI 2011, Santorini, Greece, September 18–21, 2011 (Lecture Notes in Computer Science, Vol. 6960)*, Yiannis Cotronis, Anthony Danalis, Dimitrios S. Nikolopoulos, and Jack Dongarra (Eds.). Springer, Berlin, Heidelberg, 345–355. https://doi.org/10.1007/978-3-642-24449-0_45
 - [28] Kuang-Hua Huang and Jacob A. Abraham. 1984. Algorithm-Based Fault Tolerance for Matrix Operations. *IEEE Trans. Comput.* C-33, 6 (1984), 518–528. <https://doi.org/10.1109/TC.1984.1676475>
 - [29] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP '23)*. ACM, 611–626. <https://doi.org/10.1145/3600006.3613165>
 - [30] Yannis Labrou and Tim Finin. 1994. A Semantics Approach for KQML — A General Purpose Communication Language for Software Agents. In *Proceedings of the Third International Conference on Information and Knowledge Management (CIKM '94)*. ACM, Gaithersburg, Maryland, USA, 447–455. <https://doi.org/10.1145/191246.191320>
 - [31] Richard E. Ladner and Michael J. Fischer. 1980. Parallel prefix computation. *J. ACM* 27, 4 (Oct. 1980), 831–838. <https://doi.org/10.1145/322217.322232>
 - [32] Ignacio Laguna, Ryan Marshall, Kathryn Mohror, Martin Ruefenacht, Anthony Skjellum, and Nawrin Sultana. 2019. A Large-Scale Study of MPI Usage in Open-Source HPC Applications. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '19)*. ACM, New York, NY, Article 31, 14 pages. <https://doi.org/10.1145/3295500.3356176>
 - [33] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics* 12 (2024), 157–173. https://doi.org/10.1162/tacl_a_00638
 - [34] Message Passing Interface Forum. 1994. *MPI: A Message-Passing Interface Standard*. Technical Report CS-94-230. University of Tennessee, Knoxville.
 - [35] Message Passing Interface Forum. 2023. *MPI: A Message-Passing Interface Standard, Version 4.1*. Technical Report. University of Tennessee, Knoxville. <https://www.mpi-forum.org/docs/mpi-4.1/mpi41-report.pdf>
 - [36] Microsoft. 2025. Durable Functions Overview. <https://learn.microsoft.com/en-us/azure/azure-functions/durable/durable-functions-overview>
 - [37] Model Context Protocol Contributors. 2026. Model Context Protocol Specification, revision 2026-07-28. Agentic AI Foundation / Linux Foundation. <https://modelcontextprotocol.io/specification/2026-07-28/basic/index>
 - [38] Rolf Rabenseifner. 2004. Optimization of collective reduction operations. In *Computational Science — ICCS 2004 (Lecture Notes in Computer Science, Vol. 3036)*. Springer, 1–9. https://doi.org/10.1007/978-3-540-24685-5_1
 - [39] Rolf Rabenseifner and Jesper Larsson Träff. 2004. More efficient reduction algorithms for non-power-of-two number of processors in message-passing parallel systems. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface (EuroPVM/MPI 2004) (Lecture Notes in Computer Science, Vol. 3241)*. Springer, 36–46. https://doi.org/10.1007/978-3-540-30218-6_13
 - [40] Reid G. Smith. 1980. The Contract Net Protocol: High-Level Communication and Control in a Distributed Problem Solver. *IEEE Trans. Comput.* C-29, 12 (Dec. 1980), 1104–1113. <https://doi.org/10.1109/TC.1980.1675516>
 - [41] Marc Snir, Steve Otto, Steven Huss-Lederman, David Walker, and Jack Dongarra. 1996. *MPI: The Complete Reference*. MIT Press, Cambridge, MA.
 - [42] Rajeev Thakur, Rolf Rabenseifner, and William Gropp. 2005. Optimization of collective communication operations in MPICH. *The International Journal of High Performance Computing Applications* 19, 1 (Feb. 2005), 49–66. <https://doi.org/10.1177/1094342005051521>
 - [43] Leslie G. Valiant. 1990. A Bridging Model for Parallel Computation. *Commun. ACM* 33, 8 (Aug. 1990), 103–111. <https://doi.org/10.1145/79173.79181>
 - [44] David W. Walker. 1992. *Standards for Message-Passing in a Distributed Memory Environment*. Technical Report ORNL/TM-12147. Oak Ridge National Laboratory, Oak Ridge, TN.
 - [45] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryen W. White, Doug Burger, and Chi Wang. 2023. AutoGen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. arXiv:2308.08155 [cs.AI] <https://arxiv.org/abs/2308.08155>